

```

1  %% Machine Learning Online Class - Exercise 1: Linear Regression
2
3  % Instructions
4  % -----
5  %
6  % This file contains code that helps you get started on the
7  % linear exercise. You will need to complete the following functions
8  % in this exercise:
9  %
10 %     warmUpExercise.m
11 %     plotData.m
12 %     gradientDescent.m
13 %     computeCost.m
14 %     gradientDescentMulti.m
15 %     computeCostMulti.m
16 %     featureNormalize.m
17 %     normalEqn.m
18 %
19 % For this exercise, you will not need to change any code in this file,
20 % or any other files other than those mentioned above.
21 %
22 % x refers to the population size in 10,000s
23 % y refers to the profit in $10,000s
24 %
25
26 %% Initialization
27 clear ; close all; clc
28
29 %% ===== Part 1: Basic Function =====
30 % Complete warmUpExercise.m
31 fprintf('Running warmUpExercise ... \n');
32 fprintf('5x5 Identity Matrix: \n');
33 warmUpExercise()
34
35 fprintf('Program paused. Press enter to continue.\n');
36 pause;
37
38
39 %% ===== Part 2: Plotting =====
40 fprintf('Plotting Data ...\n')
41 data = load('ex1data1.txt');
42 X = data(:, 1); y = data(:, 2);
43 m = length(y); % number of training examples
44
45 % Plot Data
46 % Note: You have to complete the code in plotData.m
47 plotData(X, y);
48
49 fprintf('Program paused. Press enter to continue.\n');
50 pause;
51
52 %% ===== Part 3: Gradient descent =====

```

```

52  % ----- Part 3: Gradient Descent -----
53  fprintf('Running Gradient Descent ...\n')
54
55  X = [ones(m, 1), data(:,1)]; % Add a column of ones to x
56  theta = zeros(2, 1); % initialize fitting parameters
57
58  % Some gradient descent settings
59  iterations = 1500;
60  alpha = 0.01;
61
62  % compute and display initial cost
63  computeCost(X, y, theta)
64
65  % run gradient descent
66  theta = gradientDescent(X, y, theta, alpha, iterations);
67
68  % print theta to screen
69  fprintf('Theta found by gradient descent: ');
70  fprintf('%f %f \n', theta(1), theta(2));
71
72  % Plot the linear fit
73  hold on; % keep previous plot visible
74  plot(X(:,2), X*theta, '-');
75  legend('Training data', 'Linear regression')
76  hold off % don't overlay any more plots on this figure
77
78  % Predict values for population sizes of 35,000 and 70,000
79  predict1 = [1, 3.5] * theta;
80  fprintf('For population = 35,000, we predict a profit of %f\n',...
81         predict1*10000);
82  predict2 = [1, 7] * theta;
83  fprintf('For population = 70,000, we predict a profit of %f\n',...
84         predict2*10000);
85
86  fprintf('Program paused. Press enter to continue.\n');
87  pause;
88
89  %% ===== Part 4: Visualizing J(theta_0, theta_1) =====
90  fprintf('Visualizing J(theta_0, theta_1) ...\n')
91
92  % Grid over which we will calculate J
93  theta0_vals = linspace(-10, 10, 100);
94  theta1_vals = linspace(-1, 4, 100);
95
96  % initialize J_vals to a matrix of 0's
97  J_vals = zeros(length(theta0_vals), length(theta1_vals));
98
99  % Fill out J_vals
100  for i = 1:length(theta0_vals)
101      for j = 1:length(theta1_vals)
102          t = [theta0_vals(i); theta1_vals(j)];
103          J_vals(i,j) = computeCost(X, y, t);

```

```
104         end
105     end
106
107
108     % Because of the way meshgrids work in the surf command, we need to
109     % transpose J_vals before calling surf, or else the axes will be flipped
110     J_vals = J_vals';
111     % Surface plot
112     figure;
113     surf(theta0_vals, theta1_vals, J_vals)
114     xlabel('\theta_0'); ylabel('\theta_1');
115
116     % Contour plot
117     figure;
118     % Plot J_vals as 15 contours spaced logarithmically between 0.01 and 100
119     contour(theta0_vals, theta1_vals, J_vals, logspace(-2, 3, 20))
120     xlabel('\theta_0'); ylabel('\theta_1');
121     hold on;
122     plot(theta(1), theta(2), 'rx', 'MarkerSize', 10, 'LineWidth', 2);
123
```